Jialin Li

jialin.li@nyu.edu | Homepage | Google Scholar

Education

New York University — Sept. 2024 – May 2029 (Expected)
Ph.D. in Computer Science, NYU Abu Dhabi Global Ph.D. Fellow — *New York, NY, USA*
Advisor: Hanan Salam

New York University Shanghai — Sept. 2019 – Jan. 2023
B.S. in Data Science — *Shanghai, China*

Research Interests

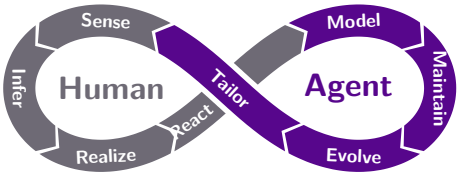

Research Focus: People-Aware LLM Agents. I build and benchmark personal agents that adapt to individual personality, context-dependent interaction preferences, and continuous user feedback.

Long-Term Goal: Co-evolving human-agent partnerships, turning daily interactions into learning signals that adapt agent harnesses to amplify human agency without constraining it.

Experience

Tencent Hunyuan, Foundation Model Team — May 2026 – Sept. 2026
Algorithm Research Intern, In-Game Companion Agents — *China*

- Built an **interactive data synthesis pipeline** with in-game state. Distilled diverse player personas at realistic distributions, enforced NPC role consistency with affective-cognitive profiles. Generated 10K+ multi-turn dialogues.
- Maintained an automated data evaluation pipeline with to ensure safety, logic while improving role-play consistencies and humor. Guided data mixing with a automated judge calibrated with human rating.
- Designed a user-centric memory framework (`L0` real-time state, `L1` recent events, `L2` single-match context, `L3` cross-match history `L4` game knowledge). Constructed training samples on memory recall and dialogue continuation.
- Performed SFT to align the bots conversational style and inject persona settings and world knowledge. Designed a **multi-objective reward model** calibrated on player behavioral signals. Further trained the bot with **offline PPO** and **online GRPO** toward multi-turn long-horizon goals, boosting composite reward scores by ∼35%.
- Built a live bad-case analysis system to process 20K sessions/day. Analyzed root cause by system layer. Closed a feedback → attribution → retraining **data flywheel**, the product served 8M total users with a peak DAU of 500K.

Publications

ACL Findings 2026 **PrefIx: Understand and Adapt to User Preference in Human-Agent Interaction**
J. Li, Z. Chen, H. Luo, H. Salam. Code

- **Problem:** LLM agents nowadays can solve tasks correctly, but their interaction still frustrates users in various ways, e.g. excessive confirmations. Existing benchmarks only evaluate task accuracy, missing the flaws in interaction.
- **Approach:** Introduced ***interaction preferences*** and **Interaction-as-a-Tool (IaaT)**. IaaT frames models interaction behaviors (confirm, disambiguate, escalate etc.) as structured tool calls, to make them evaluable with standard tool-use benchmarks. We test this across 31 preference in four dimensions (transparency, interaction pace, strategic proactiveness, robustness) with simulated users. Evaluated on seven UX dimensions by a panel of four LLM judges.
- **Findings:** Our agents match baseline accuracy while improving UX by 7.6% and alignment by 18.5%. Among the four preference dimensions, transparency was the easiest to satisfy, while robustness was the hardest.

Under Review **Why Bother Me Again? Benchmarking Personal Agent Memory for Interaction Preferences**
J. Li, Z. Chen, H. Luo, H. Salam. Code Webpage

- **Problem:** In long-term assistance, personal agents should track users' changing interaction (IX) preferences across contexts and over time to keep the interactions satisfying. However, current memory benchmarks have not been evaluating evolving IX preferences in long-horizon usage.
- **Approach:** Built an interactive environment for **on-policy memory evaluation** MemPABench. Presented four grounded simulated users with 28 interaction preferences across work and personal contexts. They used the testing

agents over 99 sessions across four months, including five preference shifts. Evaluated four memory systems across various LLM backbones on preference selection, delivered behavior, and user-perceived interaction quality.
- **Findings:** Retrieval-based memory (Mem0, MemOS, RAG) merely matches the no-memory baseline. File memory performs better, but it overgeneralizes in misapplying preferences in the wrong contexts. Worse, agents struggle with execution. Even given the right preference, they deliver the intended behavior only 59% of the time.

ACII 2025 (Oral) **Improving Personalisation in Valence and Arousal Prediction using Data Augmentation**
J. Li, M. Nwadike, H. Salam.

- **Problem:** Per-person valence and arousal models have little data for each individual.
- **Approach:** Proposed a data-driven **sample-level augmentation** that pulls similar segments from multi-person data to adapt the target model. Improved CCC on personalization benchmark without changing model architectures.

ACM MM 2025 Workshop **Personality-Aware Engagement Prediction in Online Learning**
J. Li, G. Sharma, H. Salam.

- **Problem:** Vision-based engagement prediction ignores individual differences, and existing datasets are mostly lab-collected, one-dimensional, and lack self-reports and personality.
- **Approach:** Built MOCSE, an in-the-wild dataset of Zoom lectures (72 students, 36.5 hours) with self-reported and observer engagement labels plus Big Five personality; conditioning VLMs on personality improves prediction.

arXiv 2023 **Beyond One-Size-Fits-All: A Survey of Personalized Affective Computing in Human-Agent Interaction**
J. Li, M. Elgarf, A. Waleed, H. Salam.

- First survey of personalized affective computing, with a taxonomy of data- and model-level approaches.

Skills

Evaluation & Data: Interactive benchmark and environment design, LLM-as-a-judge, multi-turn data synthesis
Post-Training: SFT, reward modeling, PPO, GRPO
Programming: Python, Rust, Shell, SQL; PyTorch, Transformers; Docker, Git, Linux, WandB

Honors & Awards

NYU Abu Dhabi Global Ph.D. Fellowship, New York University 2024 –
University Honors Scholar, New York University 2023
Dean's List of Academic Year, New York University 2021 – 2022
Dean's Undergraduate Research Fund, New York University 2020 – 2021

Service

Reviewer: ACM MM 2025, SuSS 2025, ACL 2026, ICLR 2026, IEEE TAC, EMNLP 2026
Outreach: Association for the Advancement of Affective Computing (AAAC) DEI Committee, 2025